


Binary task: which runs actually tell the two classes apart?



The test set is class-balanced, so accuracy is the midpoint of the two recalls and every point on the

dashed line scores 50%. Only Gemini leaves the bottom-right corner; Qwen3-VL sits in it, answering "Novice" almost always.